

THE AI BULLETIN

"Your daily AI briefing — Signal only, no noise" **BY ADARSH NAGAR**

MONDAY, JUNE 1, 2026

AI INTELLIGENCE BRIEFING

20 STORIES · LIVE NOW

TODAY 20

THIS WEEK 20

VOICES 18

REFRESH

Search today's stories...

01

Differences Between Opus 4.7 and Opus 4.8 on MineBench

REDDIT R/SINGULARITY · 15H AGO

02

Mystery company accidentally blew \$500 million on Claude AI in a single month

REDDIT R/ARTIFICIAL · 2D AGO

03

Anthropic Tops OpenAI to Become the World’s Most Valuable A.I. Start-Up

REDDIT R/ARTIFICIAL · 2D AGO

THE BIG PICTURE4 stories

BUILT & SHIPPED4 stories

THE SHIFT1 stories

COST GUARDRAILS MATTER

Differences Between Opus 4.7 and Opus 4.8 on MineBench

Some Notes: Average Inference Time: 24.8 min (1,487seconds) Total Cost (for 15 builds): \$41.52 Much cheaper than Opus 4.7 was, despite having the same API pricing The CoT / thinking times have clearly been streamlined (similar to what OpenAI has been doing with their latest releases) which lowers overall cost, but despite that, the...

REDDIT R/SINGULARITY · May 31 · 15h ago · GENERAL AI

\$500M BET PLACED ON THIS

Mystery company accidentally blew \$500 million on Claude AI in a single month

A mysterious, unnamed company is reported to have accidentally spent half a billion dollars in a single month on Claude AI after forgetting to set usage limits for Claude licenses for employees.

REDDIT R/ARTIFICIAL · May 30 · 2d ago · DEVELOPER TOOLS

PATTERN FORMING

\$65B BET PLACED ON THIS

Anthropic Tops OpenAI to Become the World’s Most Valuable A.I. Start-Up

Anthropic raised \$65 billion in new fund-raising that put its value at \$900 billion, ahead of OpenAI’s last valuation of \$730 billion, as the companies duel for A.I. dominance. Anthropic, once the lesser-known artificial intelligence competitor to OpenAI, has been on an inexorable rise over the past few months. The San...

REDDIT R/ARTIFICIAL · May 30 · 2d ago · DEVELOPER TOOLS

CHIP POWER SHIFT

mudler/Qwen3.6-35B-A3B-Claude-4.7-Opus-Reasoning-Distilled-APEX-MTP-GGUF just released !

Description of the module: I host 30+ free APEX MoE quantizations as independent research. My only local hardware is an NVIDIA DGX Spark (122 GB unified memory) — enough for...

SHIFT UNDERWAY

Why do the output layer weights become word vectors in Word2Vec? [D]

I'm trying to understand the intuition behind Word2Vec training using a neural network. In Word2Vec (CBOW or Skip-gram), we often hear that the weight matrices learned during training contain the vector representations (embeddings) of words. However, I don't understand why the weights of the hidden-to-output layer (or...

REDDIT R/MACHINELEARNING · May 30 · 2d ago · DEVELOPER TOOLS

What AI will search for sleeping pills and cheapest my question always gets no responses any suggestions?

REDDIT R/ARTIFICIAL · May 31 · 22h ago · GENERAL AI

AGENTS STEPPING INTO REAL WORK

Use HTML as the primary chat language of your LLM's so they can make interactive content

REDDIT R/LOCALLLAMA · Jun 1 · 7h ago · GENERAL AI

REDDIT R/LOCALLLAMA · May 31 · Yesterday · ROBOTICS & HARDWARE

ROBOTICS CROSSES A THRESHOLD

Open-weights VLA hits 80%+ task progress on 4 of 17 real-robot tasks with zero fine-tuning. Demo reel attached

REDDIT R/SINGULARITY · May 31 · Yesterday · DEVELOPER TOOLS

CHIP POWER SHIFT

I ported NVIDIA Parakeet (speech-to-text) to ggml: same output as NeMo, faster, GGUF-quantized, no Python

REDDIT R/LOCALLLAMA · May 31 · 13h ago · DEVELOPER TOOLS

AI IN THE WILD

DEVELOPER TOOLS

My home data center

System 1: Threadripper 3960x 24c 4x 3090 ti 128gb ddr4 System 2: Xeon 8352 36c 4x 5070 ti 128gb ddr4 System 3: Intel 14700k 24c 64gb ddr5 5090...

Reddit r/LocalLLaMA

GENERAL AI

GPU Prices. Buy now, or buy later?

If the Community could sound off on this, I'd be grateful. Do you think GPU prices are going to stop skyrocketing? Is this FOMO and hype driving the...

Reddit r/LocalLLaMA

CREATIVE AI

Added an old 2070 Super to my rig and I can't go back worse, now I need more

Context: I built a new system last year November before everything went to shit. I spent like 5k for a 5090, 9800X3D and 96GB RAM. Recently (last 2-3...

Reddit r/LocalLLaMA

HEALTHCARE

What I learned building a debugger for PyTorch training loops and how it changed how I think about failure diagnosis [D]

Hey r/ML, I spent the last few months building tool that hooks into PyTorch training loops to automatically detect and localize failures...

Reddit r/MachineLearning

- ALL
- RESEARCH
- BUILDING
- BUSINESS
- REAL WORLD

JUST DROPPED NEW MODELS, TOOLS, AND PRODUCTS THAT SHIPPED

1 DISPATCHES

Arabic ASR model struggling to converge during training [D]

i'm trying to train an ASR model using the LibriSpeech recipe from SpeechBrain (without the language model) on a 100-hour dataset of dialectal Arabic speech. the model architecture uses a...

Reddit r/MachineLearning · May 31 · 13h ago

News


DEVELOPER TOOLS

SHIFTING GROUND WHAT DEVS ARE BUILDING AND HOW THE INDUSTRY IS CHANGING

4 DISPATCHES

Built an AI Accelerator and opensourced it. [P]

How would you model this "strand" clustering problem? [P]



There is a huge gap in open source AI accelerators, so I implemented mine. Popular and well known ones are already legacy and doesn't support...

Reddit r/MachineLearning · May 31 · 21h ago

NewsGENERAL AI



https://preview.redd.it/llqlupnwng4h1.png?width=2188&format=png&auto=webp&s=7fae5860babaffa1c8bfdcb1468b374eb38ac5...

Reddit r/MachineLearning · May 31 · 22h ago

NewsDEVELOPER TOOLS

Bayesian Opt. GPs vs Linear models and Neural Networks for parameter optimizations [R]

Hi, Relatively new to deep learning. I wanted some opinions on which of these approaches might be best for time series data and spectral analysis. I currently use a GP and it works pretty well, b...

Reddit r/MachineLearning · May 31 · Yesterday

NewsDEVELOPER TOOLS



nvidia/Qwen3.6-35B-A3B-NVFP4 · Hugging Face

The NVIDIA Qwen3.6-35B-A3B-NVFP4 model is the quantized version of Alibaba's Qwen3.6-35B-A3B model, which is an auto-regressive language model that uses an...

Reddit r/LocalLLaMA · May 30 · Yesterday

NewsGENERAL AI

 SHIFTING GROUND HOW THE GLOBAL LANDSCAPE AND REGULATIONS ARE SHIFTING

2 DISPATCHES



Get you some GPUs, it's not worth the hacks around lack of RAM

https://preview.redd.it/w356ddr8ak4h1.png?width=550&format=png&auto=webp&s=f04238bf0d44f6defe58698c75f08d6c2581d4c...

Reddit r/LocalLLaMA · Jun 1 · 10h ago

NewsGENERAL AI



<Think> toggle button for llama.cp web chat for QWEN3.6

https://preview.redd.it/od6suf6j7g4h1.png?width=619&format=png&auto=webp&s=d31fb903ea68f58e3a641bfd275d59eeb5cce4...

Reddit r/LocalLLaMA · May 31 · 23h ago

NewsGENERAL AI

THE AI BULLETIN

"Your daily AI briefing — Signal only, no noise" **BY ADARSH NAGAR**

MONDAY, JUNE 1, 2026

AI INTELLIGENCE BRIEFING

20 STORIES · LIVE NOW

TODAY 20

THIS WEEK 20

VOICES 18

REFRESH

Search sources...

Voices — What People Are Saying

Live posts from AI researchers, builders, and communities

- All
- Reddit
- Twitter / X
- YouTube

Two Minute Papers

@TwoMinutePapers

YOUTUBE

Feynman vs. Einstein vs. Newton: Who Wins? 💬 🤖 #physics #ai #science #feynman #research

10h ago

Watch on YouTube →

Andrej Karpathy

@karpathy

LLM INTUITION

Building LLMs from scratch is the absolute best way to gain intuition. Do not just import PyTorch and call train, actually write the backprop loops! Understanding standard matrix multiplications deconstructs model scaling laws.

↑ 14,200 · 💬 2,890

View on X →

X/TWITTER

Sam Altman

@sama

AGENT REASONING

X/TWITTER

We are scaling compute structures to push frontier agent reasoning benchmarks. The next wave of agentic tools will transition from pure conversational vibe-coding to rigorous sandbox code execution environments.

↑ 9,800 · 🗨️ 1,240

View on X →

Clement Delangue

@ClementDelworker

HUGGING FACE

X/TWITTER

Follow Clement Delangue for insights on Hugging Face. RSSHub feed currently unavailable — visit their profile directly.

2d ago

View on X →

Yann LeCun

@ylecun

AI RESEARCH

X/TWITTER

Follow Yann LeCun for insights on AI Research. RSSHub feed currently unavailable — visit their profile directly.

2d ago

View on X →

Andrej Karpathy

@karpathy

LLM INTUITION

X/TWITTER

Follow Andrej Karpathy for insights on LLM Intuition. RSSHub feed currently unavailable — visit their profile directly.

2d ago

View on X →

Varun Mayya

X/TWITTER

Varun Mayya

varun-mayya

WORKFLOW AUTOMATION

X/TWITTER

Developers in India are increasingly automating legacy enterprise workflows using n8n and LangGraph multi-agent orchestrations. The developer bottleneck isn't syntax writing anymore, it's modular system architecture, prompt engineering, and context window routing.

↑ 3,450 · 🗨️ 580

View Source →

Sam Altman

@sama

OPENAI

X/TWITTER

Follow Sam Altman for insights on OpenAI. RSSHub feed currently unavailable — visit their profile directly.

2d ago

View on X →

Jim Fan

@DrJimFan

EMBODIED AI

X/TWITTER

Follow Jim Fan for insights on Embodied AI. RSSHub feed currently unavailable — visit their profile directly.

2d ago

View on X →

AK

@_akhaliq

PAPER DROPS

X/TWITTER

Follow AK for insights on Paper Drops. RSSHub feed currently unavailable — visit their profile directly.

2d ago

View on X →

Vaibhav Sisinty

X/TWITTER

THE AI BULLETIN

"Your daily AI briefing — Signal only, no noise" BY ADARSH NAGAR

MONDAY, JUNE 1, 2026

AI INTELLIGENCE BRIEFING

20 STORIES · LIVE NOW

TODAY 20

THIS WEEK 20

VOICES 18

REFRESH

Search timeline...



THIS WEEK IN AI

Weekly chronological intelligence timeline & weekly roundup

SUNDAY, MAY 31

13 STORIES TRACKED

TOP STORY

WHY IT MATTERS · COST GUARDRAILS MATTER

Differences Between Opus 4.7 and Opus 4.8 on MineBench

Some Notes: Average Inference Time: 24.8 min (1,487seconds) Total Cost (for 15 builds): \$41.52 Much cheaper than Opus 4.7 was, despite having the same API pricing The CoT / thinking times have clearly been streamlined (similar to what OpenAI has been doing with their latest releases) which lowers overall cost, but despite that, the output seems better than Opus 4.7, so that's good This is, in my opinion, one of the first Claude models in a long time that actually feels like a genuinely impressive release; its builds are actually of similar quality to GPT 5.5, though a bit more inconsistent Dur...

Reddit r/singularity • General AI

MODERATE SIGNAL

mudler/Qwen3.6-35B-A3B-Claude-4.7-Opus-Reasoning-Distilled-APEX-MTP-GGUF just released !

REDDIT R/LOCALLAMA MODERATE SIGNAL

Open-weights VLA hits 80%+ task progress on 4 of 17 real-robot tasks with zero fine-tuning. Demo reel attached

REDDIT R/SINGULARITY MODERATE SIGNAL

I ported NVIDIA Parakeet (speech-to-text) to ggml: same output as NeMo, faster, GGUF-quantized, no Python

REDDIT R/LOCALLAMA MODERATE SIGNAL

My home data center	REDDIT R/LOCALLLAMA	MODERATE SIGNAL
What AI will search for sleeping pills and cheapest my question always gets no responses any suggestions?	REDDIT R/ARTIFICIAL	MODERATE SIGNAL
GPU Prices. Buy now, or buy later?	REDDIT R/LOCALLLAMA	MODERATE SIGNAL
Arabic ASR model struggling to converge during training [D]	REDDIT R/MACHINELEARNING	MODERATE SIGNAL
Added an old 2070 Super to my rig and I can't go back worse, now I need more	REDDIT R/LOCALLLAMA	MODERATE SIGNAL
Built an AI Accelerator and opensourced it. [P]	REDDIT R/MACHINELEARNING	MODERATE SIGNAL
How would you model this "strand" clustering problem? [P]	REDDIT R/MACHINELEARNING	MODERATE SIGNAL
<Think> toggle button for llama.cp web chat for QWEN3.6	REDDIT R/LOCALLLAMA	MODERATE SIGNAL
Bayesian Opt. GPs vs Linear models and Neural Networks for parameter optimizations [R]	REDDIT R/MACHINELEARNING	MODERATE SIGNAL

SATURDAY, MAY 30

5 STORIES TRACKED

⚡ TOP STORY

WHY IT MATTERS · \$500M BET PLACED ON THIS

Mystery company accidentally blew \$500 million on Claude AI in a single month

A mysterious, unnamed company is reported to have accidentally spent half a billion dollars in a single month on Claude AI after forgetting to set usage limits for Claude licenses for employees.

Reddit r/artificial • Developer Tools

MODERATE SIGNAL

Anthropic Tops OpenAI to Become the World's Most Valuable A.I. Start-Up	REDDIT R/ARTIFICIAL	MODERATE SIGNAL
nvidia/Qwen3.6-35B-A3B-NVFP4 · Hugging Face	REDDIT R/LOCALLLAMA	MODERATE SIGNAL